

INTRODUCTION TO OPENACC

João Paulo Navarro, Solutions Architect

LECTURE OUTLINE

Topics to be Covered

- What is OpenACC and Why Should You Care?
- Profile-driven Development
- First Steps with OpenACC
- Data Movement
- Hackathon

INTRODUCTION TO OPENACC

3 WAYS TO ACCELERATE APPLICATIONS

Applications

Libraries

Easy to use
Most Performance

Compiler
Directives

Easy to use
Portable code

OpenACC

Programming
Languages

Most Performance
Most Flexibility

OPENACC IS...

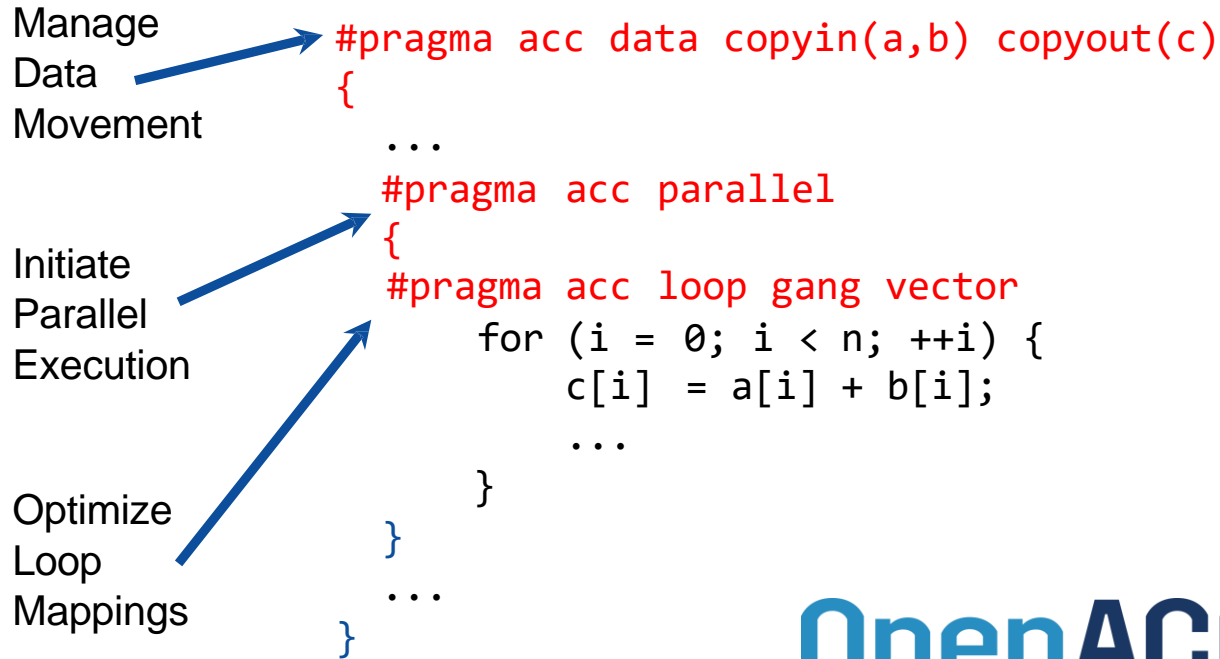
a directives-based **parallel programming model**
designed for **performance**
and **portability**.

Add Simple Compiler Directive

```
main()
{
    <serial code>
    #pragma acc kernels
    {
        <parallel code>
    }
}
```



OpenACC Directives



- Incremental
- Single source
- Interoperable
- Performance portable
- CPU, GPU, Manycore

OpenACC
Directives for Accelerators

OPENACC

Incremental

- Maintain existing sequential code
- Add annotations to expose parallelism
- After verifying correctness, annotate more of the code

Single Source

- Rebuild the same code on multiple architectures
- Compiler determines how to parallelize for the desired machine
- Sequential code is maintained

Low Learning Curve

- OpenACC is meant to be easy to use, and easy to learn
- Programmer remains in familiar C, C++, or Fortran
- No reason to learn low-level details of the hardware.

OPENACC

Incremental

- Maintain existing sequential code
- Add annotations to expose parallelism
- After verifying correctness, annotate more of the code

Enhance Sequential Code

```
#pragma acc parallel loop
for( i = 0; i < N; i++ )
{
    < loop code >
}

#pragma acc parallel loop
for( i = 0; i < N; i++ )
{
    < loop code >
}
```

Begin with a working sequential code.

Parallelize it with OpenACC.

Rerun the code to verify correctness and performance

OPENACC

Supported Platforms

POWER
Sunway
x86 CPU
AMD GPU
NVIDIA GPU
PEZY-SC

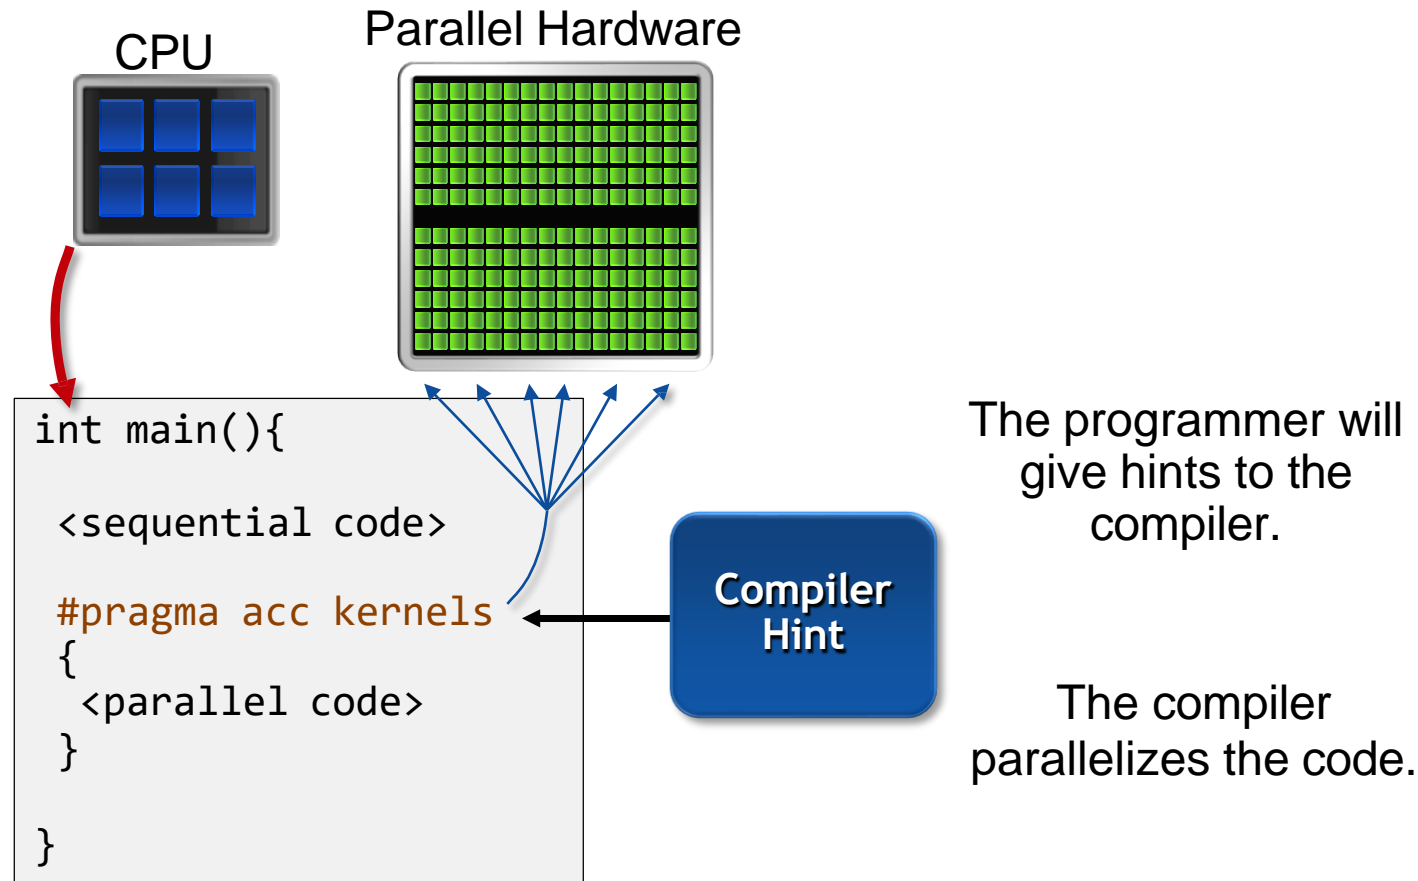
Single Source

- Rebuild the same code on multiple architectures
- Compiler determines how to parallelize for the desired machine
- Sequential code is maintained

The compiler can **ignore** your OpenACC code additions, so the same code can be used for **parallel** or **sequential** execution.

```
int main(){  
  
...  
  
    #pragma acc parallel loop  
    for(int i = 0; i < N; i++)  
        < loop code >  
  
}
```

OPENACC



Low Learning Curve

- OpenACC is meant to be easy to use, and easy to learn
- Programmer remains in familiar C, C++, or Fortran
- No reason to learn low-level details of the hardware.

OPENACC

Incremental

- Maintain existing sequential code
- Add annotations to expose parallelism
- After verifying correctness, annotate more of the code

Single Source

- Rebuild the same code on multiple architectures
- Compiler determines how to parallelize for the desired machine
- Sequential code is maintained

Low Learning Curve

- OpenACC is meant to be easy to use, and easy to learn
- Programmer remains in familiar C, C++, or Fortran
- No reason to learn low-level details of the hardware.

DIRECTIVE-BASED HPC PROGRAMMING

Who's Using OpenACC

3 OF TOP 5 HPC APPS



5 OF 13 CAAR CODES



2 OF LAST 9 FINALISTS



450 DOMAIN EXPERTS



ACCELERATED APPS



100,000 DOWNLOADS



GAUSSIAN 16



Mike Friesch, Ph.D.
President and
CEO,
Gaussian, Inc.

“Using OpenACC allowed us to continue development of our fundamental algorithms and software capabilities simultaneously with the GPU-related work. In the end, we could use the same code base for SMP, cluster/ network and GPU parallelism. PGI's compilers were essential to the success of our efforts.”




Image courtesy: ANSYS

ANSYS FLUENT



Sunil Sathya
Lead Software Developer
ANSYS Fluent

“We've effectively used OpenACC for heterogeneous computing in ANSYS Fluent with impressive performance. We're now applying this work to more of our models and new platforms.”



VASP



Prof. Georg Kresse
Computational Materials Physics
University of Vienna

“For VASP, OpenACC is the way forward for GPU acceleration. Performance is similar and in some cases better than CUDA C, and OpenACC dramatically decreases GPU development and maintenance efforts. We're excited to collaborate with NVIDIA and PGI as an early adopter of CUDA Unified Memory.”



COSMO



Dr. Oliver Fuhrer
Senior Scientist
Matscorder

“OpenACC made it practical to develop for GPU-based hardware while retaining a single source for almost all the COSMO physics code.”



E3SM



Mark A. Taylor
Multiphysics Applications
Sandia

“The CAAR project provided us with early access to Summit hardware and access to PGI compiler experts. Both of these were critical to our success. PGI's OpenACC support remains the best available and is competitive with much more intrusive programming model approaches.”




Image courtesy: Oak Ridge National Laboratory



NUMECA FINE/Open



David Gubeloff
Lead Software Developer
NUMECA

“Porting our unstructured C++ CFD solver FINE/Open to GPUs using OpenACC would have been impossible two or three years ago, but OpenACC has developed enough that we're now getting some really good results.”



SYNOPSIS



Dr. Lutz Schneider
Senior R&D Engineer
Synopsis Inc.

“Using OpenACC, we've GPU-accelerated the Synopsis TCAD Sentaurus Device EMW simulator to speed up optical simulations of image sensors. GPUs are key to improving simulation throughput in the design of advanced image sensors.”




Image courtesy: NCAR

MPAS-A



Richard Loft
Director, Technology Development
NCAR

“Our team has been evaluating OpenACC as a pathway to performance portability for the Model for Prediction (MPAS) atmospheric model. Using this approach on the MPAS dynamical core, we have achieved performance on a single P100 GPU equivalent to 2-7 dual socketed Intel Xeon nodes on our new Cheyenne supercomputer.”





VMD



John Stone
Senior Research Programmer
Buckham Institute
University of Illinois

“Due to Amdahl's law, we need to port more parts of our code to the GPU if we're going to speed it up. But the sheer number of routines poses a challenge. OpenACC directives give us a low-cost approach to getting at least some speed-up out of these second-tier routines. In many cases it's completely sufficient because with the current algorithms, GPU performance is bandwidth-bound.”



GTC



Zhihong Lin
Professor and Principal Investigator
UC Irvine

“Using OpenACC our scientists were able to achieve the acceleration needed for integrated fusion simulation with a minimum investment of time and effort in learning to program GPUs.”



OpenACC
More Science. Less Programming

GAMERA



Takuma Yamaguchi, Kohji Fujita, Shojiro Ichimaru, Masaki Hirai, Lutz Schneider
The University of Tokyo

“With OpenACC and a compute node based on NVIDIA's Tesla P100 GPU, we achieved more than a 14X speed up over a K Computer node running our earthquake disaster simulation code.”





Map courtesy: University of Tokyo

SANJEEVINI



Abhilash Jayaram
Project Scientist
Indian Institute of Technology
New Delhi

“In an academic environment maintenance and speedup of existing codes is a tedious task. OpenACC provides a great platform for computational scientists to accomplish both tasks without involving a lot of efforts or manpower in speeding up the entire computational task.”




IBM-CFD



Sreenath Roy
Assistant Professor
Mechanical Engineering Department
Indian Institute of Technology Kharagpur

“OpenACC can prove to be a handy tool for computational engineers and researchers to obtain fast solution of non-linear dynamics problems. In incompressible boundary incompressible CFD, we have obtained order of magnitude reduction in computing time by porting several components of our legacy codes to GPU. Especially the routines involving search algorithm and matrix solvers have been well-accelerated to improve the overall scalability of the code.”




PWscf (Quantum ESPRESSO)



Filippo Spina
Senior Contributor
Quantum ESPRESSO group

“CUDA Fortran gives us the full performance potential of the CUDA programming model and NVIDIA GPUs. While leveraging the potential of explicit data movement, ISCUF KERNELS directives give us productivity and source code maintainability. It's the best of both worlds.”



MAS



Ronald M. Caplan
Computational Scientist
Predictive Science Inc.

“Adding OpenACC into MAS has given us the ability to migrate medium-sized simulations from a multi node CPU cluster to a single multi-GPU server. The implementation yielded a portable single-source code for both CPU and GPU runs. Future work will add OpenACC to the remaining model features, enabling GPU accelerated realistic solar storm modeling.”



OPENACC SYNTAX

OPENACC SYNTAX

Syntax for using OpenACC directives in code

C/C++

```
#pragma acc directive clauses  
<code>
```

Fortran

```
!$acc directive clauses  
<code>
```

- A ***pragma*** in C/C++ gives instructions to the compiler on how to compile the code. Compilers that do not understand a particular pragma can freely ignore it.
- A ***directive*** in Fortran is a specially formatted comment that likewise instructs the compiler in its compilation of the code and can be freely ignored.
- “***acc***” informs the compiler that what will come is an OpenACC directive
- ***Directives*** are commands in OpenACC for altering our code.
- ***Clauses*** are specifiers or additions to directives.

EXAMPLE CODE

LAPLACE HEAT TRANSFER

Introduction to lab code - visual


We will observe a simple simulation of heat distributing across a metal plate.

We will apply a consistent heat to the top of the plate.

Then, we will simulate the heat distributing across the plate.


Very Hot

Room Temp



EXAMPLE: JACOBI ITERATION

- Iteratively converges to correct value (e.g. Temperature), by computing new values at each point from the average of neighboring points.
- Common, useful algorithm
- Example: Solve Laplace equation in 2D: $\nabla^2 f(x, y) = 0$



$$A_{k+1}(i, j) = \frac{A_k(i-1, j) + A_k(i+1, j) + A_k(i, j-1) + A_k(i, j+1)}{4}$$

JACOBI ITERATION: C CODE

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
  
    iter++;  
}
```



Iterate until converged



Iterate across matrix
elements



Calculate new value from
neighbors



Compute max error for
convergence




Swap input/output arrays

PROFILE-DRIVEN DEVELOPMENT

OPENACC DEVELOPMENT CYCLE

- **Analyze** your code to determine most likely places needing parallelization or optimization.
- **Parallelize** your code by starting with the most time consuming parts and check for correctness.
- **Optimize** your code to improve observed speed-up from parallelization.



PROFILING SEQUENTIAL CODE

Profile Your Code

Obtain detailed information about how the code ran.


This can include information such as:

- Total runtime
- Runtime of individual routines
- Hardware counters

Identify the portions of code that took the longest to run. We want to focus on these “hotspots” when parallelizing.

Lab Code: Laplace Heat Transfer

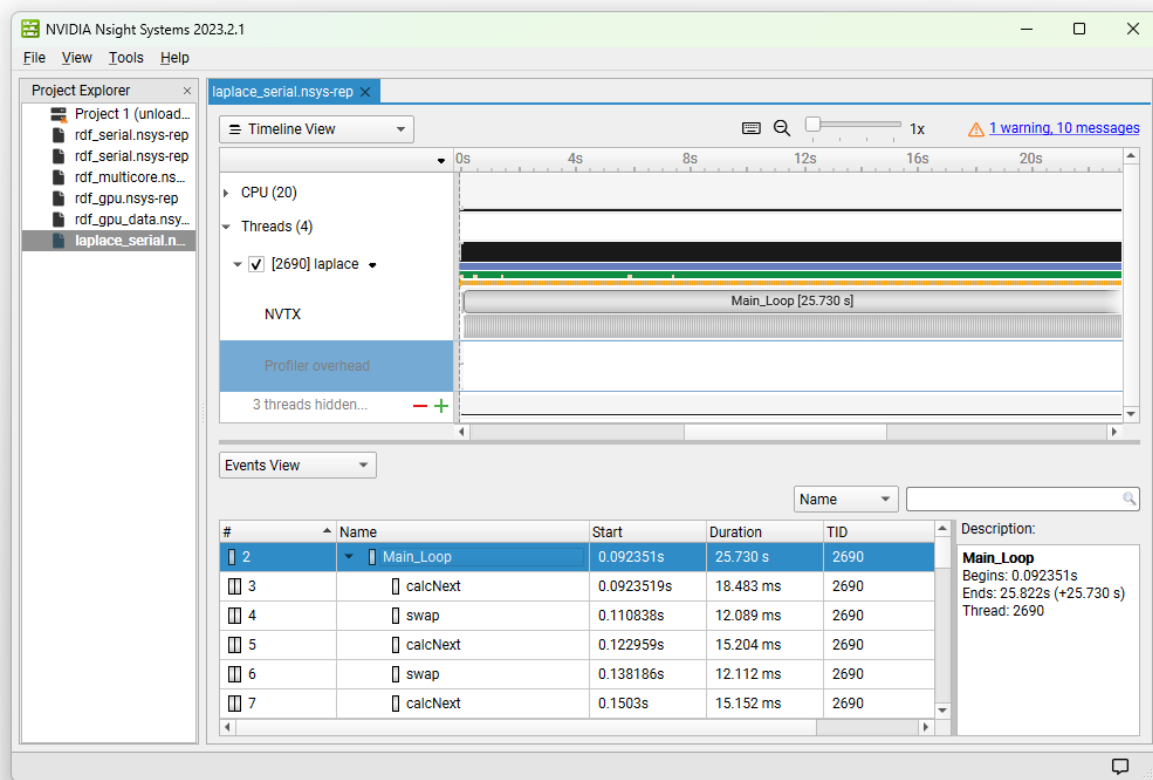
Total Runtime: 39.43 seconds



PROFILING SEQUENTIAL CODE

First sight when using NSIGHT SYSTEMS

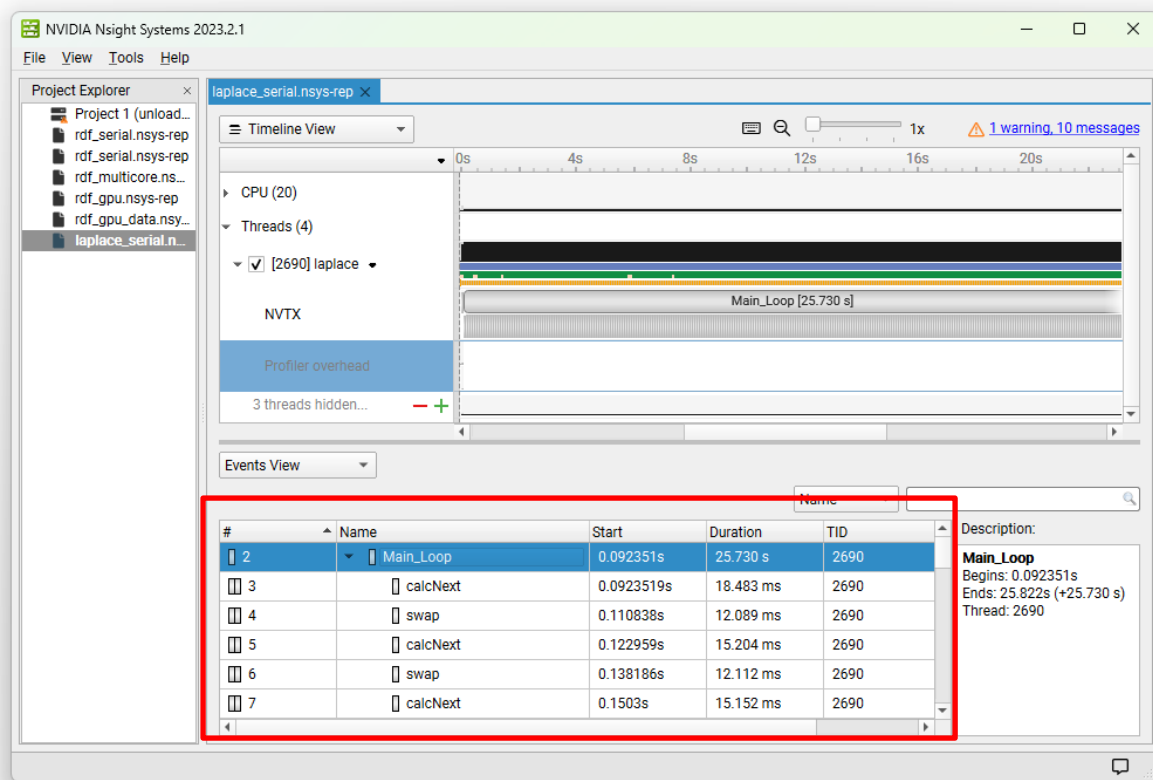
- Profiling a simple, sequential code
- Our sequential program will run on the CPU
- To view information about how our code ran, we add NVTx markers in our code!



PROFILING SEQUENTIAL CODE

CPU Details

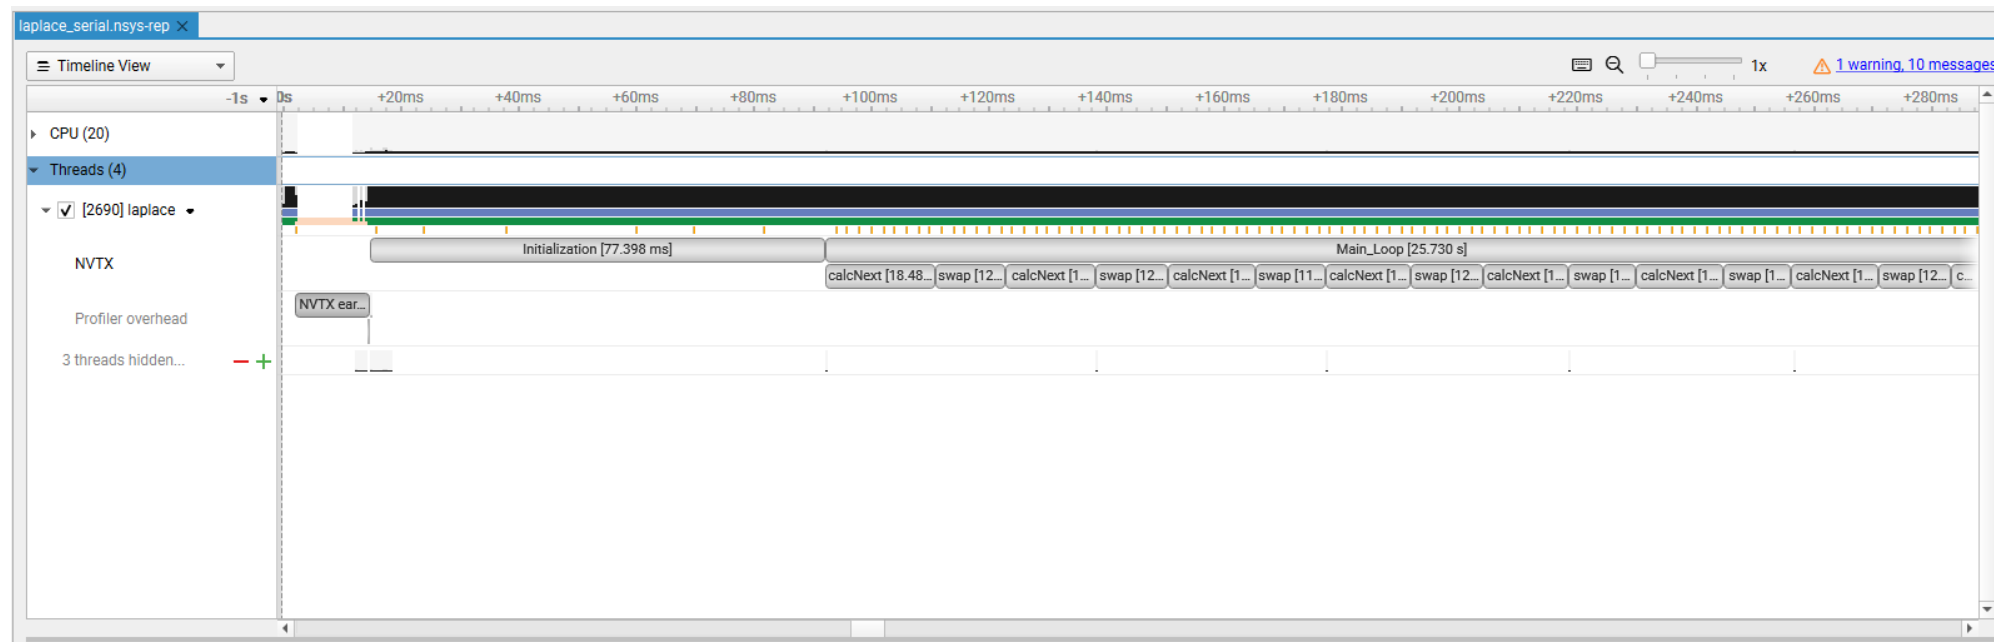
- On NVTX line we see the total time spent at each markup
- Double click in NVTX line, will open the event tab
- We will expand this information, and see more details about our code



PROFILING SEQUENTIAL CODE

CPU Details

- We can zoom in to get a close into the time spent in each NVTX marker.



OPENACC PARALLEL LOOP DIRECTIVE


OPENACC PARALLEL DIRECTIVE

Expressing parallelism

```
#pragma acc parallel  
{
```

When encountering the ***parallel*** directive, the compiler will generate *1 or more parallel gangs*, which execute redundantly.

```
}
```



OPENACC PARALLEL DIRECTIVE


Expressing parallelism

```
#pragma acc parallel  
{
```

```
    loop  
    for(int i = 0; i < N; i++)  
    {  
        // Do Something  
    }
```

```
}
```

This loop will be
executed redundantly
on each gang



OPENACC PARALLEL DIRECTIVE


Expressing parallelism

```
#pragma acc parallel  
{
```

```
  for(int i = 0; i < N; i++)  
  {  
    // Do Something  
  }
```

```
}
```

This means that each
gang will execute the
entire loop



OPENACC PARALLEL DIRECTIVE

Parallelizing a single loop

C/C++

```
#pragma acc parallel
{
    #pragma acc loop
    for(int i = 0; i < N; i++)
        a[i] = 0;
}
```

Fortran

```
!$acc parallel
!$acc loop
do i = 1, N
    a(i) = 0
end do
!$acc end parallel
```

- Use a **parallel** directive to mark a region of code where you want parallel execution to occur
- This parallel region is marked by curly braces in C/C++ or a start and end directive in Fortran
- The **loop** directive is used to instruct the compiler to parallelize the iterations of the next loop to run across the parallel gangs

OPENACC PARALLEL DIRECTIVE

Parallelizing a single loop

C/C++

```
#pragma acc parallel loop
for(int i = 0; i < N; i++)
    a[i] = 0;
```

Fortran

```
!$acc parallel loop
do i = 1, N
    a(i) = 0
end do
```

- This pattern is so common that you can do all of this in a single line of code
- In this example, the parallel loop directive applies to the next loop
- This directive both marks the region for parallel execution and distributes the iterations of the loop.
- When applied to a loop with a data dependency, parallel loop may produce incorrect results

OPENACC PARALLEL DIRECTIVE


Expressing parallelism

```
#pragma acc parallel  
{
```

```
    #pragma acc loop  
    for(int i = 0; i < N; i++)  
    {  
        // Do Something  
    }
```

```
}
```

The **loop** directive informs the compiler which loops to parallelize.



OPENACC PARALLEL LOOP DIRECTIVE

Parallelizing many loops

```
#pragma acc parallel loop
for(int i = 0; i < N; i++)
    a[i] = 0;

#pragma acc parallel loop
for(int j = 0; j < M; j++)
    b[j] = 0;
```

- To parallelize multiple loops, each loop should be accompanied by a parallel directive
- Each parallel loop can have different loop boundaries and loop optimizations
- Each parallel loop can be parallelized in a different way
- This is the recommended way to parallelize multiple loops. Attempting to parallelize multiple loops within the same parallel region may give performance issues or unexpected results

PARALLELIZE WITH OPENACC PARALLEL LOOP

```
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```

Parallelize first loop nest,
max *reduction* required.

Parallelize second loop.

We didn't detail *how* to
parallelize the loops, just *which*
loops to parallelize.

REDUCTION CLAUSE

- The **reduction** clause takes many values and “reduces” them to a single value, such as in a sum or maximum
- Each thread calculates its part
- The compiler will perform a final reduction to produce a **single global result** using the specified operation

```
for( i = 0; i < size; i++ )  
    for( j = 0; j < size; j++ )  
        for( k = 0; k < size; k++ )  
            c[i][j] += a[i][k] * b[k][j];
```

```
for( i = 0; i < size; i++ )  
    for( j = 0; j < size; j++ )  
        double tmp = 0.0f;  
        #pragma parallel acc loop \  
            reduction(+:tmp)  
        for( k = 0; k < size; k++ )  
            tmp += a[i][k] * b[k][j];  
        c[i][j] = tmp;
```

REDUCTION CLAUSE OPERATORS

Operator	Description	Example
+	Addition/Summation	<code>reduction(+:sum)</code>
*	Multiplication/Product	<code>reduction(*:product)</code>
max	Maximum value	<code>reduction(max:maximum)</code>
min	Minimum value	<code>reduction(min:minimum)</code>
&	Bitwise and	<code>reduction(&:val)</code>
 	Bitwise or	<code>reduction(:val)</code>
&&	Logical and	<code>reduction(&&:val)</code>
 	Logical or	<code>reduction(:val)</code>

BUILD AND RUN THE CODE



NVIDIA HPC SDK

PROGRAMMING THE NVIDIA PLATFORM

CPU, GPU, and Network

ACCELERATED STANDARD LANGUAGES

ISO C++, ISO Fortran

```
std::transform(par, x, x+n, y, y,  
    [=] (float x, float y){ return y +  
    a*x; }  
);
```

```
do concurrent (i = 1:n)  
    y(i) = y(i) + a*x(i)  
enddo
```

```
import cunumeric as np  
...  
def saxpy(a, x, y):  
    y[:] += a*x
```

INCREMENTAL PORTABLE OPTIMIZATION

OpenACC, OpenMP

```
#pragma acc data copy(x,y) {  
...  
std::transform(par, x, x+n, y, y,  
    [=] (float x, float y){  
        return y + a*x;  
    });  
...  
}  
  
#pragma omp target data map(x,y) {  
...  
std::transform(par, x, x+n, y, y,  
    [=] (float x, float y){  
        return y + a*x;  
    });  
...  
}
```

PLATFORM SPECIALIZATION

CUDA

```
__global__  
void saxpy(int n, float a,  
    float *x, float *y) {  
    int i = blockIdx.x*blockDim.x +  
        threadIdx.x;  
    if (i < n) y[i] += a*x[i];  
}  
  
int main(void) {  
    ...  
    cudaMemcpy(d_x, x, ...);  
    cudaMemcpy(d_y, y, ...);  
  
    saxpy<<<(N+255)/256,256>>>(...);  
  
    cudaMemcpy(y, d_y, ...);  
}
```

ACCELERATION LIBRARIES

Core

Math

Communication

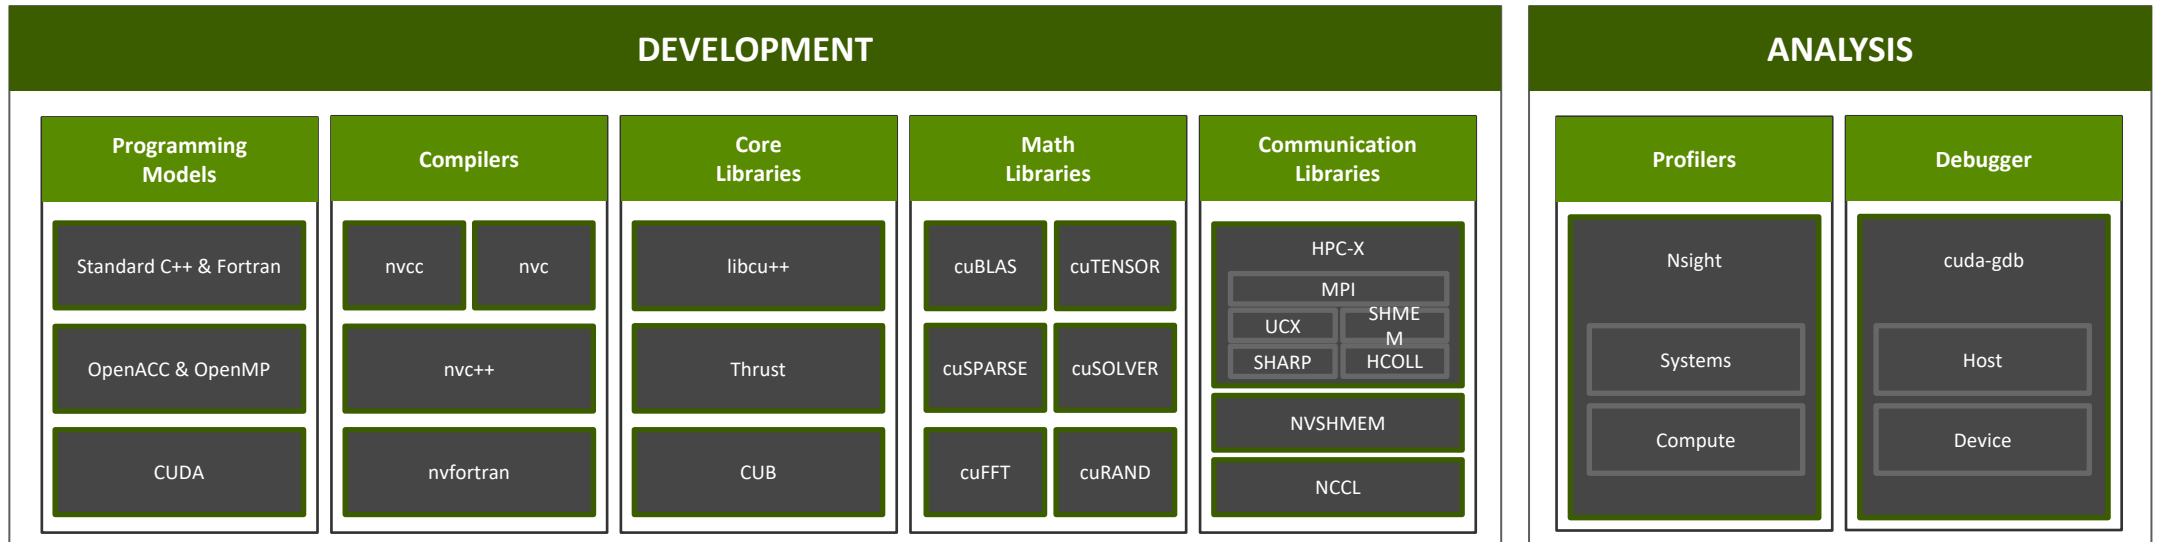
Data Analytics

AI

Quantum

NVIDIA HPC SDK

Available at developer.nvidia.com/hpc-sdk, on NGC, via Spack, and in the Cloud



Develop for the NVIDIA Platform: GPU, CPU and Interconnect
Libraries | Accelerated C++ and Fortran | Directives | CUDA
7-8 Releases Per Year | Freely Available

NVIDIA HPC SDK COMPILER BASICS

`nvc`, `nvc++` and `nvfortran`

- The command to compile C code is '`nvc`'
- The command to compile C++ code is '`nvc++`'
- The command to compile Fortran code is '`nvfortran`'
- The `-fast` flag instructs the compiler to optimize the code to the best of its abilities

```
$ nvc -fast main.c  
$ nvc++ -fast main.cpp  
$ nvfortran -fast main.F90
```

NVIDIA HPC SDK COMPILER BASICS

-Minfo flag

- The -Minfo flag will instruct the compiler to print feedback about the compiled code
- **-Minfo=accel** will give us information about what parts of the code were accelerated via OpenACC
- **-Minfo=opt** will give information about all code optimizations
- **-Minfo=all** will give all code feedback, whether positive or negative

```
$ nvc -fast -Minfo=all main.c  
$ nvc++ -fast -Minfo=all main.cpp  
$ nvfortran -fast -Minfo=all main.f90
```


NVIDIA HPC SDK COMPILER BASICS

-acc flag

- The -acc flag enables building OpenACC code for an Accelerator (acc)
- -acc=multicore – Build the code to run across threads on a multicore CPU
- -acc=gpu -gpu:managed – Build the code for an NVIDIA GPU and manage the data movement for me (next lecture)

```
$ nvc -fast -Minfo=accel -acc=gpu -gpu:managed main.c  
$ nvc++ -fast -Minfo=accel -acc=gpu -gpu:managed main.cpp  
$ nvfortran -fast -Minfo=accel -acc=gpu -gpu:managed main.f90
```

BUILDING THE CODE (MULTICORE)

```
$ nvc -fast -acc=multicore -Minfo=accel laplace2d.c
```

```
main:
```

```
63, Generating Multicore code
```

```
64, #pragma acc loop gang
```

```
64, Accelerator restriction: size of the GPU copy of Anew,A is unknown  
Generating reduction(max:error)
```

```
66, Loop is parallelizable
```

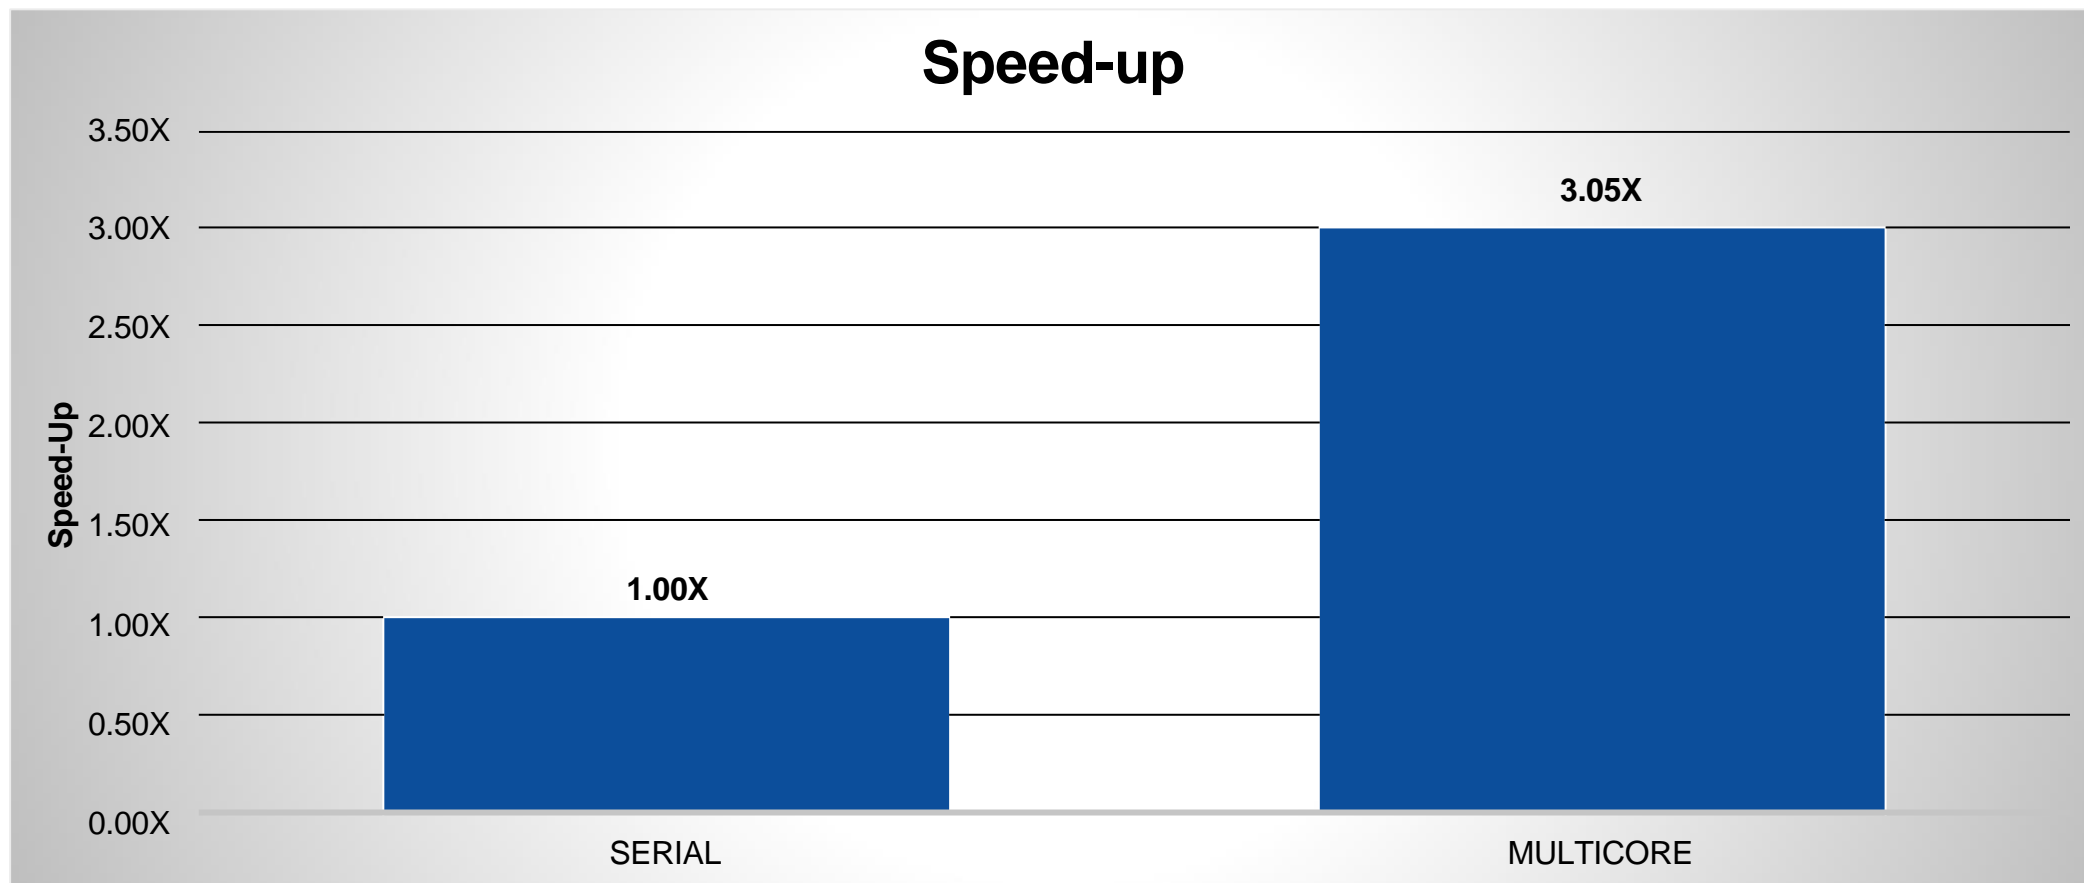
```
74, Generating Multicore code
```

```
75, #pragma acc loop gang
```

```
75, Accelerator restriction: size of the GPU copy of Anew,A is unknown
```

```
77, Loop is parallelizable
```

OPENACC SPEED-UP

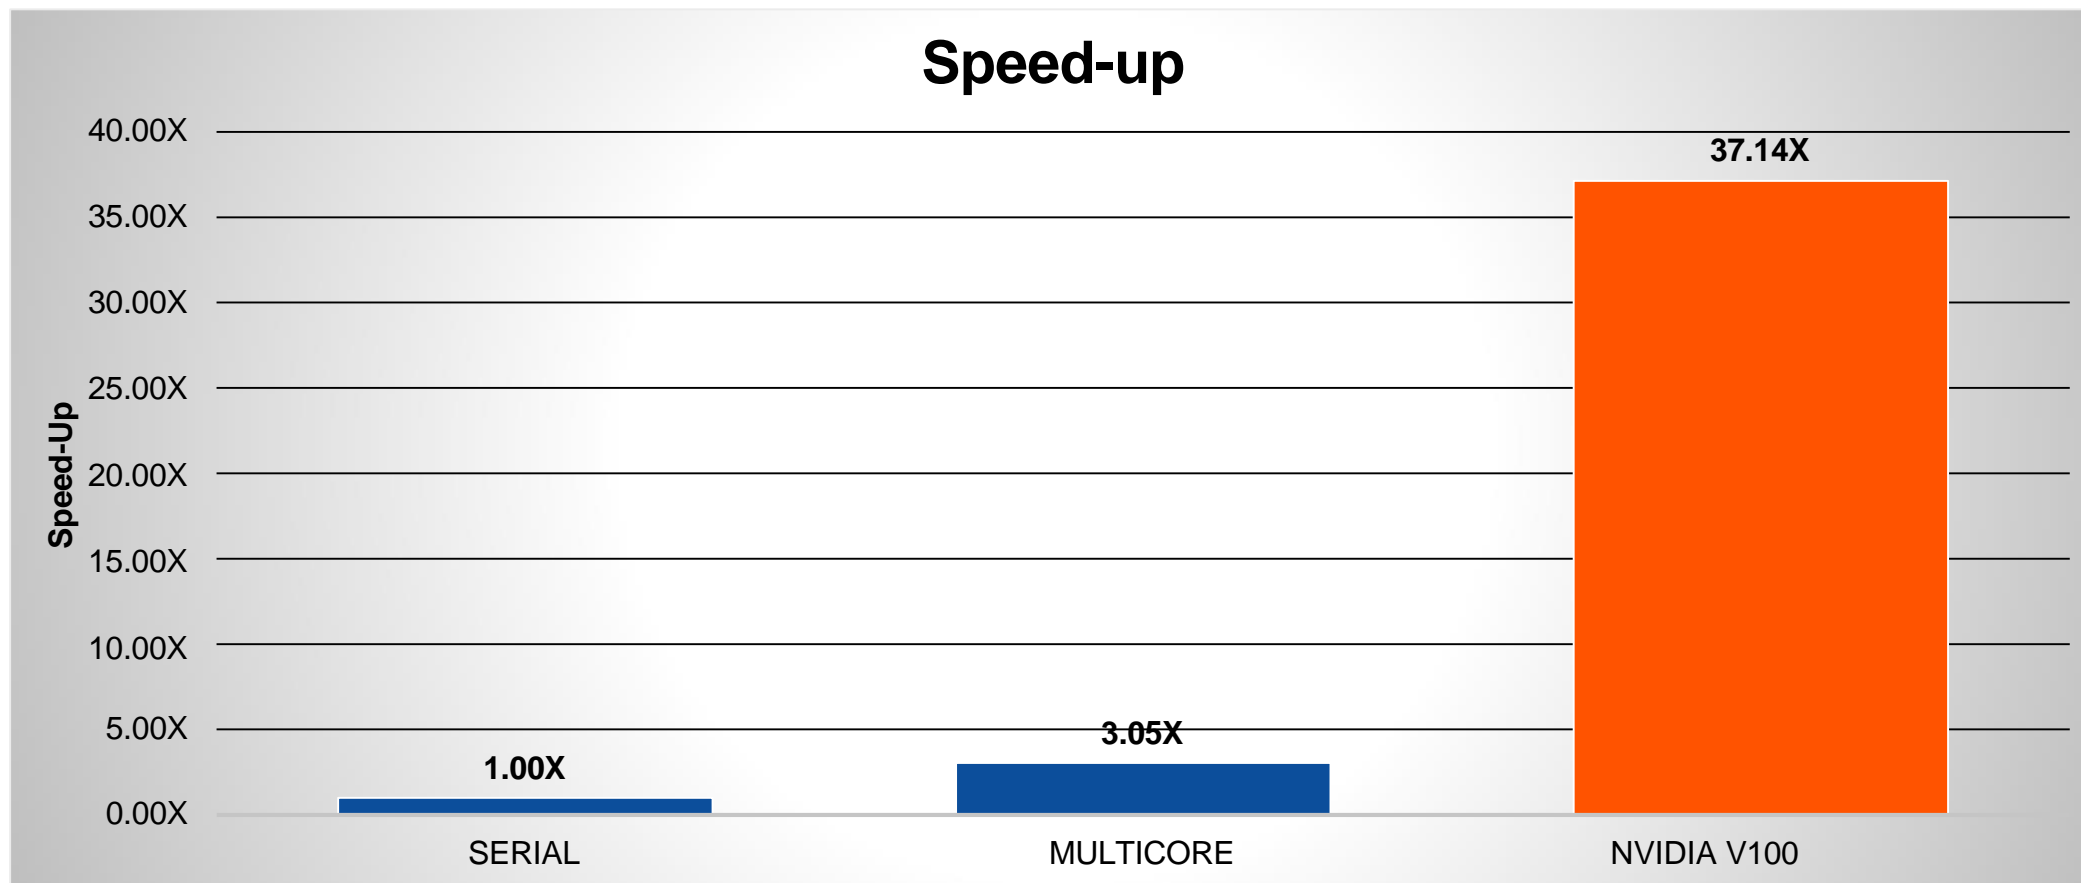


BUILDING THE CODE (GPU)

```
$ nvc -fast -acc=gpu -gpu:managed -Minfo=accel laplace2d.c main:
```

```
63, Accelerator kernel generated
Generating NVIDIA GPU code
64, #pragma acc loop gang /* blockIdx.x */
Generating reduction(max:error)
66, #pragma acc loop vector(128) /* threadIdx.x */
63, Generating implicit copyin(A[:])
Generating implicit copyout(Anew[:])
Generating implicit copy(error)
66, Loop is parallelizable
74, Accelerator kernel generated
Generating Tesla code
75, #pragma acc loop gang /* blockIdx.x */
77, #pragma acc loop vector(128) /* threadIdx.x */
74, Generating implicit copyin(Anew[:])
Generating implicit copyout(A[:])
77, Loop is parallelizable
```

OPENACC SPEED-UP



CLOSING REMARKS

KEY CONCEPTS

In this lecture we discussed...

- What is OpenACC
- How profile-driven programming helps you write better code
- How to parallelize loops using OpenACC's **parallel loop** directive to improve time to solution

Next Lecture:

- Managing your data with OpenACC

OPENACC DATA MANAGEMENT


João Paulo Navarro, Solutions Architect

CPU AND GPU MEMORIES

CPU + GPU

Physical Diagram

- CPU memory is larger, GPU memory has more bandwidth
- CPU and GPU memory are usually separate, connected by an I/O bus (traditionally PCI-e)
- Any data transferred between the CPU and GPU will be handled by the I/O Bus
- The I/O Bus is relatively slow compared to memory bandwidth
- The GPU cannot perform computation until the data is within its memory




CUDA UNIFIED MEMORY

CUDA UNIFIED MEMORY


Simplified Developer Effort

Commonly referred to as
“managed memory.”

Without Managed Memory



With Managed Memory



CUDA MANAGED MEMORY

Usefulness

- Handling explicit data transfers between the host and device (CPU and GPU) can be difficult
- The PGI compiler can utilize CUDA Managed Memory to defer data management
- This allows the developer to concentrate on parallelism and think about data movement as an optimization

```
$ nvc -fast -acc=gpu -gpu:managed -Minfo=accel main.c
```


```
$ nvfortran -fast -acc=gpu -gpu:managed -Minfo=accel main.f90
```

MANAGED MEMORY

Limitations


- The programmer will almost always be able to get better performance by manually handling data transfers
- Memory allocation/deallocation takes longer with managed memory
- Cannot transfer data asynchronously
- Currently only available from PGI on NVIDIA GPUs.

With Managed Memory



SPEC ACCEL 1.2 OPENACC BENCHMARKS

OpenACC with Unified Memory vs OpenACC Data Directives



PGI 18.4 Compilers OpenACC SPEC ACCEL™ 1.2 performance measured June, 2018
SPEC® and the benchmark name SPEC ACCEL™ are registered trademarks of the Standard Performance Evaluation Corporation.

PGI

58 NVIDIA

OpenACC
More Science. Less Programming.

NVIDIA

aws

Linux Academy

* Slide Courtesy of PGI

LAST SESSION WE USED UNIFIED MEMORY

Now let's make our code run without.

Why?


- Better data flow control and performance
- Currently the data always arrives “Just Too Late”, let's do better

BASIC DATA MANAGEMENT

BASIC DATA MANAGEMENT

Between the host and device


- The **host** is traditionally a CPU
- The **device** is some parallel accelerator
- When our target hardware is multicore, the host and device are the same, meaning that their memory is also the same
- There is no need to explicitly manage data when using a shared memory accelerator, such as the multicore target



BASIC DATA MANAGEMENT

Between the host and device

- When the target hardware is a GPU data will usually need to migrate between CPU and GPU memory
- Each array used on the GPU must be allocated on the GPU
- When data changes on the CPU or GPU the other must be updated



TRY TO BUILD WITHOUT “MANAGED”

Change the compiling line to remove “managed” part

```
nvc -acc=gpu -gpu:managed -Minfo=accel laplace2d.c jacobi.c
```

```
laplace2d.c:
```

```
PGC-S-0155-Compiler failed to translate accelerator region (see -Minfo messages): Could not find allocated-variable index for symbol (laplace2d.c: 47)
```

```
calcNext:
```

```
47, Accelerator kernel generated
```

```
Generating Tesla code
```

```
48, #pragma acc loop gang /* blockIdx.x */
```

```
Generating reduction(max:error)
```

```
50, #pragma acc loop vector(128) /* threadIdx.x */
```

```
48, Accelerator restriction: size of the GPU copy of Anew,A is unknown
```

```
50, Loop is parallelizable
```

```
PGC-F-0704-Compilation aborted due to previous errors. (laplace2d.c)
```

```
PGC/x86-64 Linux 18.7-0: compilation aborted
```

```
jacobi.c:
```

DATA SHAPING

DATA CLAUSES

`copy(list)`

Allocates memory on GPU and copies data from host to GPU when entering region and copies data to the host when exiting region.

Principal use: For many important data structures in your code, this is a logical default to input, modify and return the data.

`copyin(list)`

Allocates memory on GPU and copies data from host to GPU when entering region.

Principal use: Think of this like an array that you would use as just an input to a subroutine.

`copyout(list)`

Allocates memory on GPU and copies data to the host when exiting region.

Principal use: A result that isn't overwriting the input data structure.

`create(list)`

Allocates memory on GPU but does not copy.

Principal use: Temporary arrays.

ARRAY SHAPING

- Sometimes the compiler needs help understanding the *shape* of an array
- The first number is the start index of the array
- In C/C++, the second number is how much data is to be transferred
- In Fortran, the second number is the ending index

```
copy(array[starting_index:length])
```

C/C++

```
copy(array(starting_index:ending_index))
```

Fortran

ARRAY SHAPING (CONT.)

Multi-dimensional Array shaping

```
copy(array[0:N][0:M])
```

C/C++

Both of these examples copy a 2D array to the device

```
copy(array(1:N, 1:M))
```

Fortran

ARRAY SHAPING (CONT.)

Partial Arrays

```
copy(array[i*N/4:N/4])
```

C/C++

Both of these examples copy only $\frac{1}{4}$ of the full array

```
copy(array(i*N/4:i*N/4+N/4))
```

Fortran

OPTIMIZED DATA MOVEMENT

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;
```

```
#pragma acc parallel loop reduction(max:err) copyin(A[0:n*m]) copy(Anew[0:n*m])  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }
```

Data clauses
provide necessary
“shape” to the
arrays.


```
#pragma acc parallel loop copyin(Anew[0:n*m]) copyout(A[0:n*m])  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

TRY TO BUILD WITHOUT “MANAGED”

Change the compiling line to remove “managed” part

```
nvc -acc=gpu -Minfo=accel laplace2d.c jacobi.c
laplace2d.c:
calcNext:
    47, Generating copyin(A[:m*n]) Accelerator
       kernel generated Generating Tesla code
    48, #pragma acc loop gang /* blockIdx.x */
       Generating reduction(max:error)
    50, #pragma acc loop vector(128) /* threadIdx.x */
       47, Generating implicit copy(error)
       Generating copy(Anew[:m*n])
    50, Loop is parallelizable
swap:
    62, Generating copyin(Anew[:m*n])
       Generating copyout(A[:m*n]) Accelerator
       kernel generated Generating Tesla code
    63, #pragma acc loop gang /* blockIdx.x */
    65, #pragma acc loop vector(128) /* threadIdx.x */ 65, Loop
is parallelizable
jacobi.c:
```

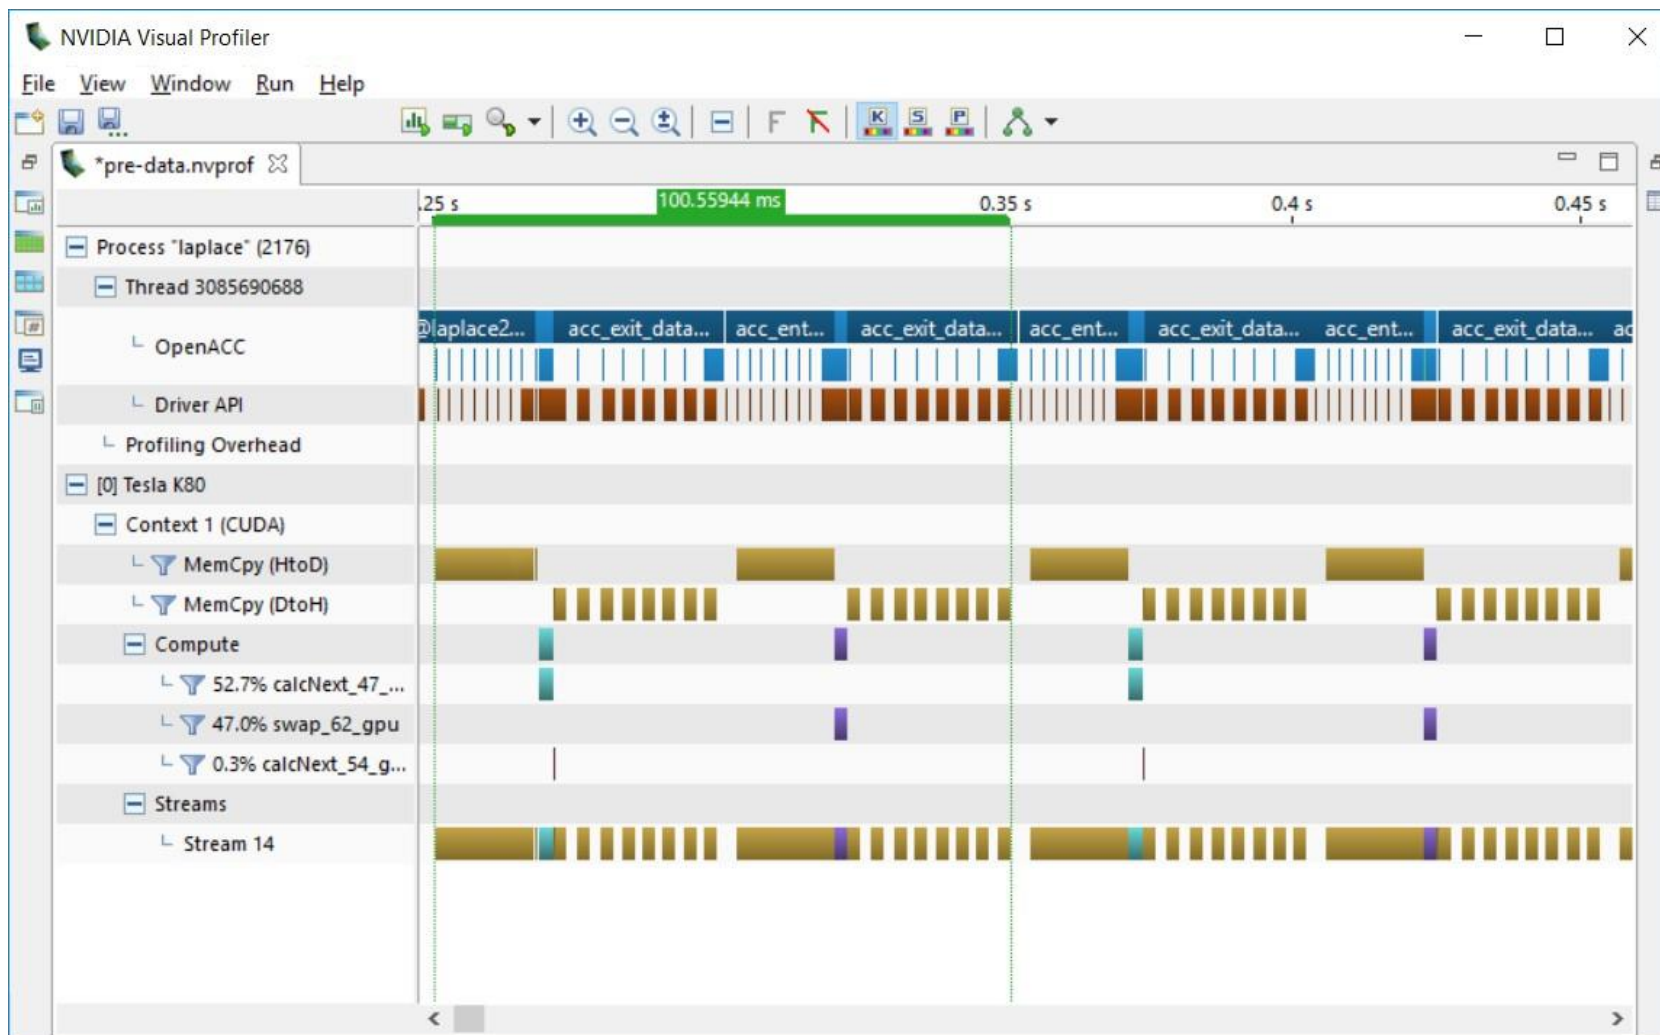
OPENACC ~~SPEED-UP~~ SLOWDOWN



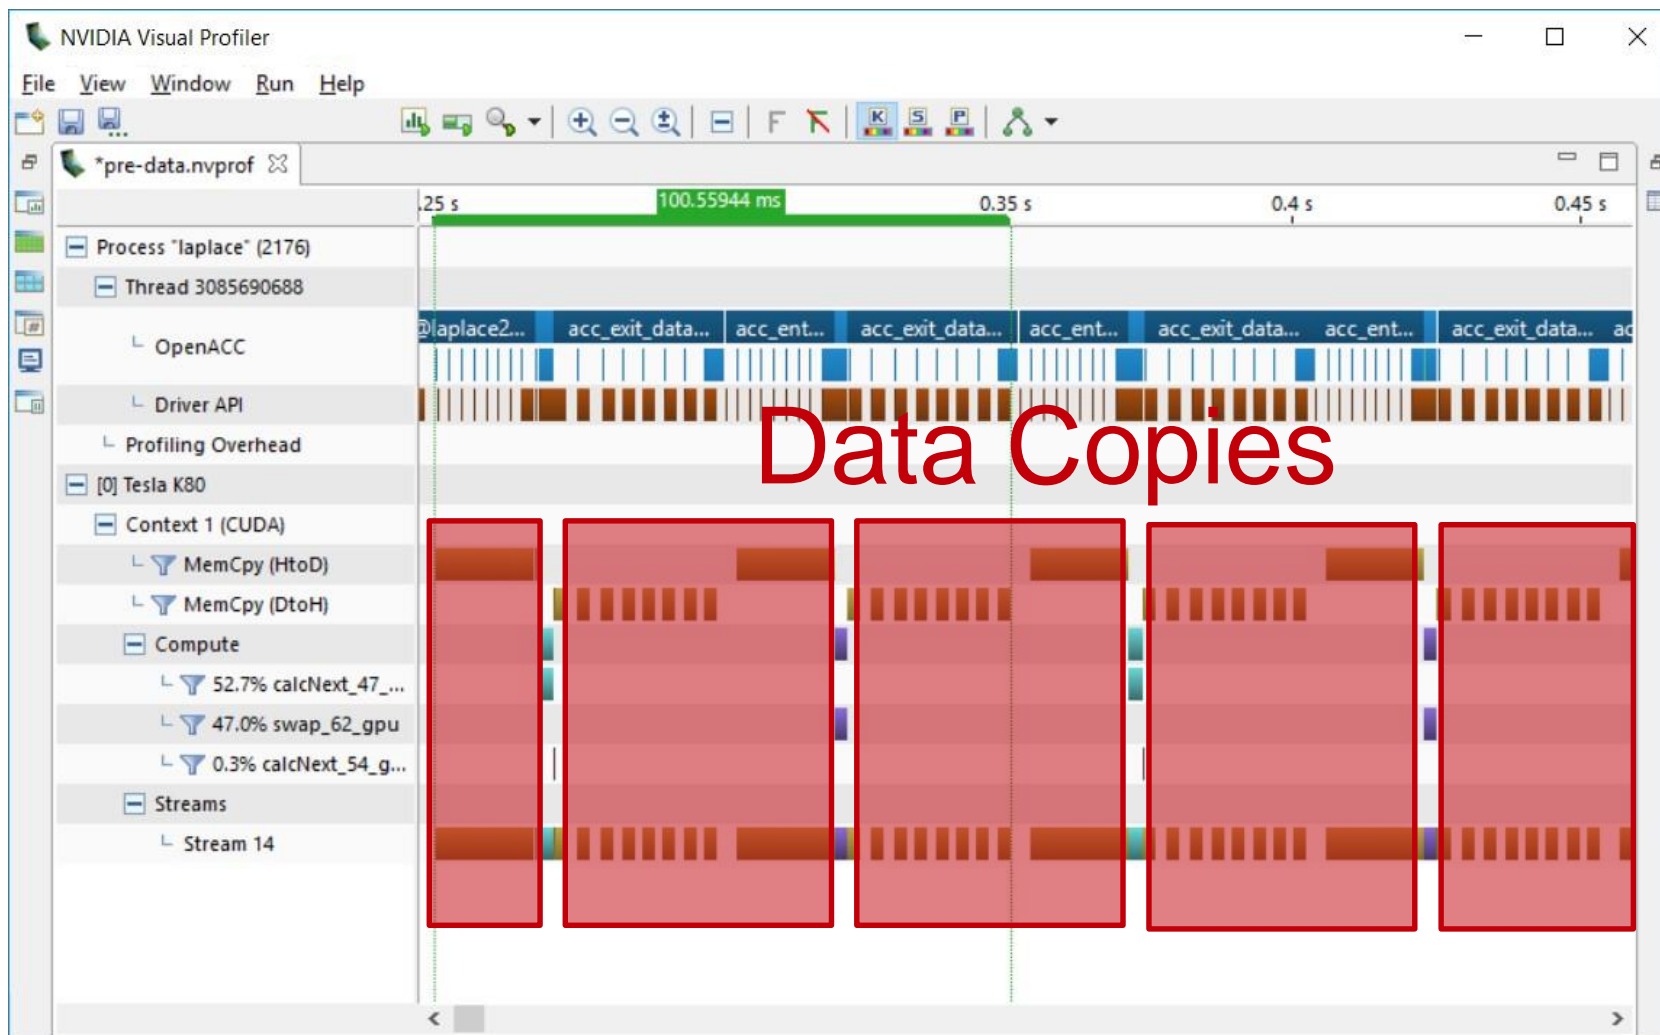
WHAT WENT WRONG?

- The code now has all of the information necessary to build without managed memory, but it runs much slower.
- Profiling tools are here to help!


APPLICATION PROFILE (2 STEPS)



APPLICATION PROFILE (2 STEPS)



RUNTIME BREAKDOWN



Nearly all of our time is spent moving data to/from the GPU

OPTIMIZED DATA MOVEMENT

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;
```

```
#pragma acc parallel loop reduction(max:err) copyin(A[0:n*m]) copy(Anew[0:n*m])
```

```
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {
```

```
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);
```

```
            err = max(err, abs(Anew[j][i] - A[j][i]));
```

```
        }  
    }
```

```
#pragma acc parallel loop copyin(Anew[0:n*m]) copyout(A[0:n*m])
```

```
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {
```

```
            A[j][i] = Anew[j][i];
```

```
        }  
    }
```

```
    iter++;
```

```
}
```

Currently we're
copying to/from the
GPU for each loop,
can we reuse it?

OPTIMIZE DATA MOVEMENT

OPENACC DATA DIRECTIVE

Definition

- The data directive defines a lifetime for data on the device beyond individual loops
- During the region data is essentially “owned by” the accelerator
- Data clauses express shape and data movement for the region

```
#pragma acc data clauses  
{  
    < Sequential and/or Parallel code >  
}
```

```
!$acc data clauses  
    < Sequential and/or Parallel code >  
!$acc end data
```

OPTIMIZED DATA MOVEMENT

```
#pragma acc data copy(A[:n*m]) copyin(Anew[:n*m])
while ( err > tol && iter < iter_max ) {
    err=0.0;
```

```
#pragma acc parallel loop reduction(max:err) copyin(A[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }
```

```
#pragma acc parallel loop copyin(Anew[0:n*m]) copyout(A[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
```

Copy A to/from the accelerator only when needed.

Copy initial condition of Anew, but not final value

REBUILD THE CODE


```
pgcc -fast -ta=tesla -Minfo=accel laplace2d_uvm.c
main:
```

```
60, Generating copy(A[:m*n])
    Generating copyin(Anew[:m*n])
64, Accelerator kernel generated
    Generating Tesla code
    64, Generating reduction(max:error)
    65, #pragma acc loop gang /* blockIdx.x */
    67, #pragma acc loop vector(128) /* threadIdx.x */
67, Loop is parallelizable
75, Accelerator kernel generated
    Generating Tesla code
    76, #pragma acc loop gang /* blockIdx.x */
    78, #pragma acc loop vector(128) /* threadIdx.x */
78, Loop is parallelizable
```



Now data movement only happens at our data region.

OPENACC SPEED-UP



WHAT WE'VE LEARNED SO FAR

- CUDA Unified (Managed) Memory is a powerful porting tool
- GPU programming without managed memory often requires data shaping
- Moving data at each loop is often inefficient
- The OpenACC Data region can decouple data movement and computation

DATA SYNCHRONIZATION

OPENACC UPDATE DIRECTIVE

update: Explicitly transfers data between the host and the device

Useful when you want to synchronize data in the middle of a data region

Clauses:

self: makes host data agree with device data

device: makes device data agree with host data

```
#pragma acc update self(x[0:count])  
#pragma acc update device(x[0:count])
```


C/C++

```
!$acc update self(x(1:end_index))  
!$acc update device(x(1:end_index))
```

Fortran

OPENACC UPDATE DIRECTIVE

```
#pragma acc update device(A[0:N])
```



The data must exist on both the CPU and device for the update directive to work.

```
#pragma acc update self(A[0:N])
```

SYNCHRONIZE DATA WITH UPDATE

```
int* A=(int*) malloc(N*sizeof(int))
#pragma acc data create(A[0:N])
while( timesteps++ < numSteps )
{
    #pragma acc parallel loop
    for(int i = 0; i < N; i++){
        a[i] *= 2;
    }

    if (timestep % 100 ) {
        #pragma acc update self(A[0:N])
        checkpointAToFile(A, N);
    }
}
```

- Sometimes data changes on the host or device inside a data region
- Ending the data region and starting a new one is expensive
- Instead, update the data so that the host and device data are the same
- Examples: File I/O, Communication, etc.

CLOSING REMARKS

KEY CONCEPTS

In this lecture we discussed...

- Differences between CPU, GPU, and Unified Memories
- OpenACC Array Shaping
- OpenACC Data Clauses
- OpenACC Structured Data Region
- OpenACC Update Directive
- OpenACC Unstructured Data Directives

Next Week: Loop Optimizations